

+ 21

# Automatically Process Your Operations in Bulk With Coroutines

FRANCESCO ZOFFOLI



20  
21



# About me

- Professional software engineer
- Passionate about C++, used throughout career
- Currently building monitoring systems at Facebook
- Author of the book “C++ Fundamentals” - Packt

makers.f.dev@gmail.com

[github.com/MakersF](https://github.com/MakersF)

# Presentation

- The Problem
- The Idea
- Coroutines in C++20
- The Implementation
- Benchmark

# Presentation

Code on github: [github.com/MakersF/cppcon-2021-corobatch](https://github.com/MakersF/cppcon-2021-corobatch)

# The Problem

# The Problem

Bulk operations tend to be better than single operations

Bulk operation: processing multiple actions or data (a batch) at once

# The Problem

- networking
- file I/O
- memory allocation
- GPU draw calls
- ...

Common property:

$\text{cost} = \text{FixedCost} + O(\# \text{operation/data})$

# The Problem

## Networking latency

round trip  $\gg$  data transfer

## Write disk

disk seek  $\gg$  data write

## DB table scan

loading data in memory  $\gg$  conditions to check per row



# What's better?

Multiple aspect

- throughput
- latency
- others: rate limit, cost, ...

Different properties are important in different situations

# The Problem

Use case:

“Send a notification to a group of users, respecting their preferences”

# The Problem

## Bulk operations

- fetching preferences
- sending notifications

## Signature

```
vector<UserPrefs> getUserPrefs(vector<UserId> userIds);  
vector<bool> sendNotifs(vector<EmailAddress> addresses);
```

## Example - No Batching

```
for(const User& user : users) {  
    const UserPrefs prefs = getUserPrefs({user.id}).at(0);  
    if (prefs.wantsEmailNotification) {  
        sendNotifs({prefs.notificationEmail});  
    }  
}
```

# The Problem

But batching is important!

## Example - Batching

```
vector<UserId> userIds;  
transform(users.begin(), users.end(),  
          back_inserter(userIds),  
          [](auto& user) { return user.id; });  
vector<UserPrefs> preferences = getUserPrefs(userIds);
```

## Example - Batching

```
remove_if(preferences.begin(), preference.end(),  
          [](auto& pref) { return !pref.wantsEmailNotification;});
```

## Example - Batching

```
vector<EmailAddress> emails;  
transform(preferences.begin(), preference.end(),  
          back_inserter(emails),  
          [](auto& pref) { return pref.notificationEmail; });  
sendNotification(emails);
```



## Example - Batching

```
vector<UserId> userIds;  
transform(users.begin(), users.end(),  
          back_inserter(userIds),  
          [](auto& user) { return user.id; });  
vector<UserPrefs> preferences = getUserPrefs(userIds);  
  
remove_if(preferences.begin(), preference.end(),  
          [](auto& pref) { return !pref.wantsEmailNotification;});  
  
vector<EmailAddress> emails;  
transform(preferences.begin(), preference.end(),  
          back_inserter(emails),  
          [](auto& pref) { return pref.notificationEmail; });  
sendNotification(emails);
```

# Readability

Code is different

- requires effort
- extra memory
- cannot use const

# Benchmark

## Simulation

- 10ms delay per call
- 50us delay per item
- 100 users
- 50 emails to send

# Benchmark

Test	Time (us)	Iterations	Baseline
No Batching	1'552'542	13	49.99
Manual Batching	31'055	492	-

A 50x difference

Questions?

# The Idea

# The Idea

Can we keep

- the readability of the first version, and
- the batching behaviour of the second?

# The Idea

What if we could gather all the parameters passed to the function?

```
for(const User& user : users) {  
    const UserPrefs prefs = getUserPrefs({user.id}).at(0);  
    if (prefs.wantsEmailNotification) {  
        sendNotifs({prefs.notificationEmail});  
    }  
}
```



# The Idea

Imagine we have 4 users

```
for(const User& user : users) {  
    const UserPrefs prefs = getUserPrefs({user.id}).at(0);  
    if (prefs.wantsEmailNotification) {  
        sendNotifs({prefs.notificationEmail});  
    }  
}
```

# The Idea

```
for(user : users) {  
  prefs = getUserPrefs(user.id);  
  if (prefs.notify) {  
    sendNotification(prefs.email);  
  }  
}
```

User 1

```
prefs = getUserPrefs(user.id);  
if (prefs.notify) {  
  sendNotification(prefs.email);  
}
```

User 2

```
prefs = getUserPrefs(user.id);  
if (prefs.notify) {  
  sendNotification(prefs.email);  
}
```

User 3

```
prefs = getUserPrefs(user.id);  
if (prefs.notify) {  
  sendNotification(prefs.email);  
}
```

...

User n

```
prefs = getUserPrefs(user.id);  
if (prefs.notify) {  
  sendNotification(prefs.email);  
}
```

# The Idea

User 1


```
➡ prefs = getUserPrefs(user.id);  
  if (prefs.notify) {  
    sendNotification(prefs.email);  
  }
```

User 2


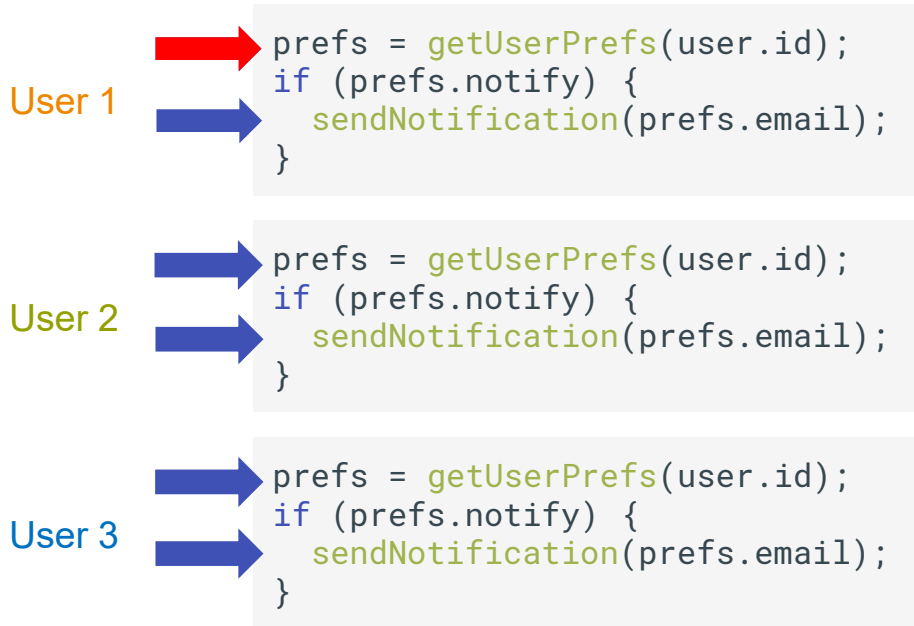
```
➡ prefs = getUserPrefs(user.id);  
  if (prefs.notify) {  
    sendNotification(prefs.email);  
  }
```

User 3

```
➡ prefs = getUserPrefs(user.id);  
  if (prefs.notify) {  
    sendNotification(prefs.email);  
  }
```



# The Idea



# The Idea

1. Execute until a batch operation
2. Record the argument, and stop
3. If not enough arguments, go to point 1
4. Execute the batch operation
5. Resume the stopped executions with the results

# The Idea

What do we need?

- Suspend execution
- Resume execution

Coroutines allow us to do exactly that


Questions?

# Coroutines



# Coroutines

## Resumable functions



# Coroutines in C++

How does a coroutine look like?

```
task my_coro() {  
    cout << "foo begin\n";  
    co_await async_operation();  
    co_return;  
}
```

# Coroutines in C++

To be a coroutine a function must

- use one of: `co_await`, `co_yield` and `co_return`
- return a type which satisfies the requirements for coroutines

# Coroutines in C++ - promise

Requirements for the return type

- Must contain a nested type called `promise_type`
- `return_type::promise_type` must implement several functions

# Coroutines in C++ - promise

The body is rewritten

```
{
    promise_type promise;
    co_await promise.initial_suspend();
    try {
        [function-body]
    } catch (...) {
        promise.unhandled_exception();
    }
    final_suspend:
    co_await promise.final_suspend();
}
```

# Coroutines in C++ - promise

## Let's build a valid coroutine return type

```
struct task {
    struct promise_type {

    };
};
```

# Coroutines in C++ - promise

When `my_coro` is invoked, the compiler

1. creates the coroutine state
2. instantiates `task::promise_type`
3. calls `get_return_object`

```
struct task {  
    struct promise_type {  
        task get_return_object();  
  
};  
};
```

# Coroutines in C++ - promise

It calls `initial_suspend()` and `co_await`s the result

```
struct task {  
    struct promise_type {  
        task get_return_object();  
        Awaitable initial_suspend();  
  
    };  
};
```



# Coroutines in C++ - co\_await

## What happens in `co_await` awaitable?

The compiler inserts calls to `awaitable` methods

```
struct Awaitable {  
  
};
```

# Coroutines in C++ - co\_await

First, it calls `await_ready()`

`true` → immediately resumes the coroutine

`false` → suspends the coroutine

```
struct Awaitable {  
    bool await_ready();  
  
};
```

# Coroutines in C++ - co\_await

After suspending the coroutine, it calls `await_suspend()`

Should schedule the coroutine to be resumed

```
struct Awaitable {  
    bool await_ready();  
    auto await_suspend(std::coroutine_handle<>);  
};
```

# Coroutines in C++ - co\_await

Once resumed, it calls `await_resume()`

The `co_await` expression evaluates to the returned object

```
struct Awaitable {  
    bool await_ready();  
    auto await_suspend(std::coroutine_handle<>);  
    T await_resume();  
};
```

# Coroutines in C++ - handle

What is the `std::coroutine_handle` we saw a few times?

# Coroutines in C++ - handle

A handle

- a not owning pointer
- controls the coroutine

# Coroutines in C++ - handle

`std::coroutine_handle<>` methods

- `resume()`: resume the associated coroutine
- `destroy()`: destroy the associated coroutine

```
struct coroutine_handle<> {  
    void resume();  
    void destroy();  
  
};
```

# Coroutines in C++ - handle

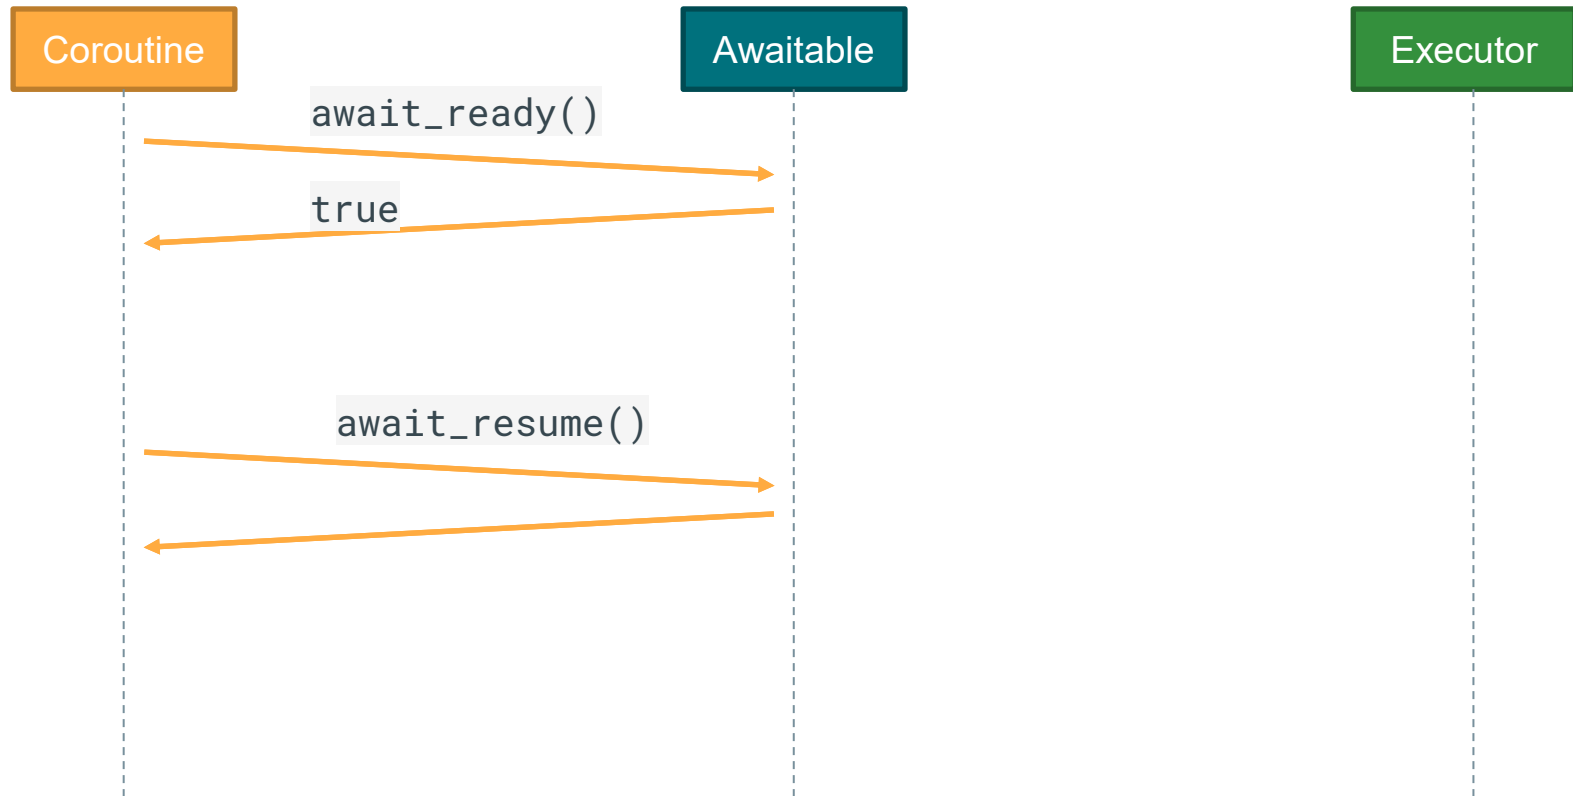
`std::coroutine_handle<promise_type>` additional methods

- `promise()`: returns the associated promise
- `from_promise()`: creates a handle from the promise

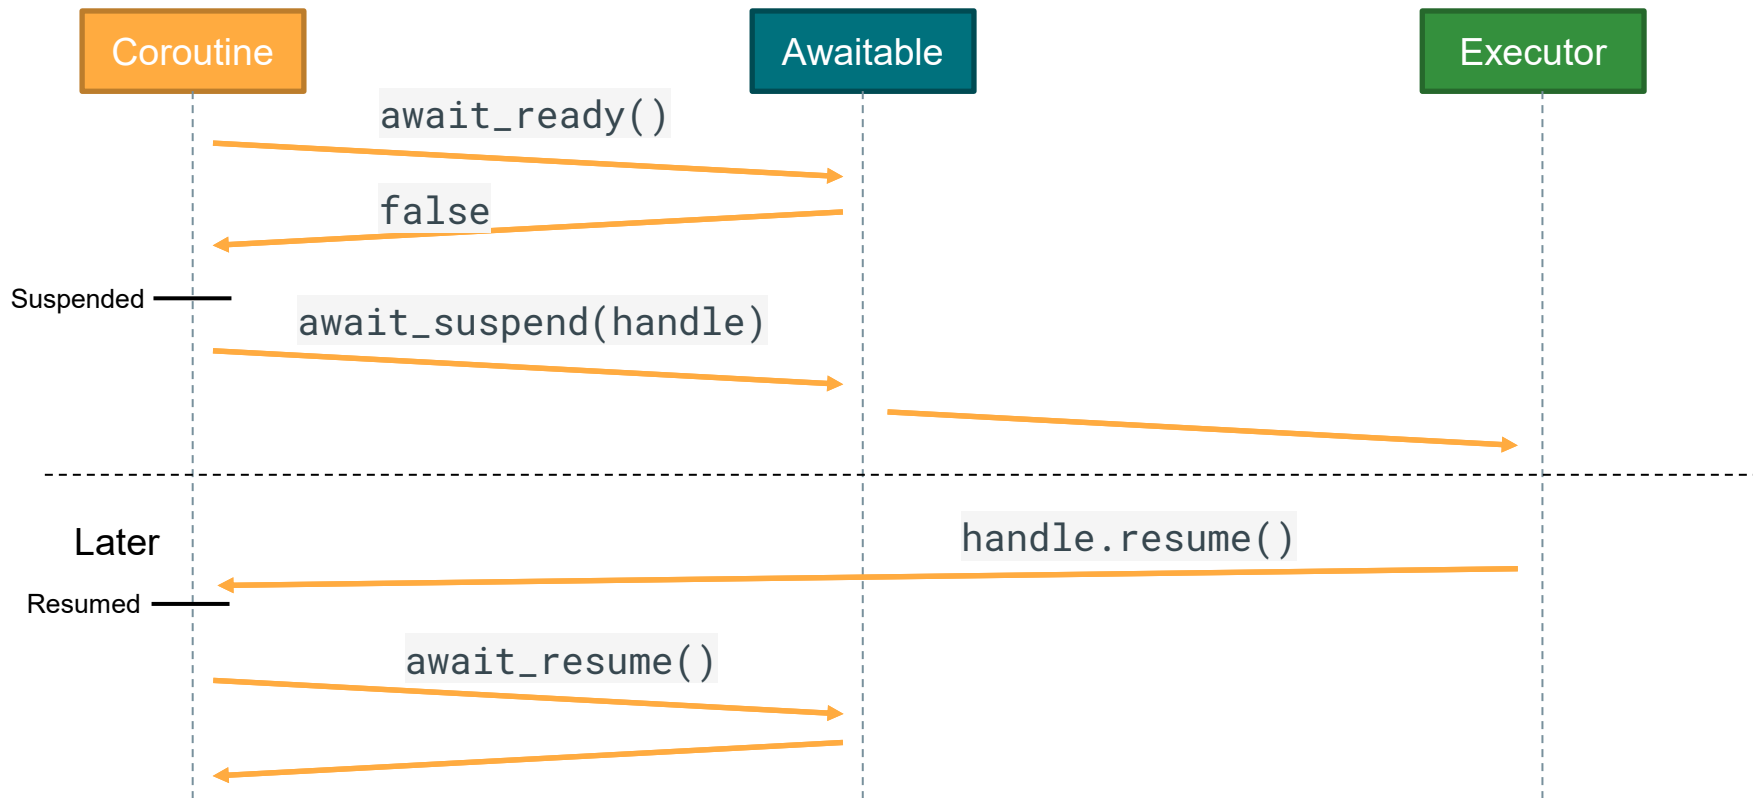
```
struct coroutine_handle<promise_type> {  
    void resume();  
    void destroy();  
    promise_type& promise();  
    static coroutine_handle<promise_type> from_promise(promise_type&);  
};
```



# Coroutines in C++ - co\_await



# Coroutines in C++ - co\_await



# Coroutines in C++ - promise

Back to the promise type!

```
struct task {  
    struct promise_type {  
        task get_return_object();  
        Awaitable initial_suspend();  
  
    };  
};
```

# Coroutines in C++ - promise

How to return values

# Coroutines in C++ - promise

`co_return expr; → return_value(expr)`

```
struct task {  
    struct promise_type {  
        task get_return_object();  
        Awaitable initial_suspend();  
        void return_value(T);  
    };  
};
```

# Coroutines in C++ - promise

`co_return;` → `return_void()`

```
struct task {  
    struct promise_type {  
        task get_return_object();  
        Awaitable initial_suspend();  
        void return_void();  
    };  
};
```

# Coroutines in C++ - promise

Uncaught exceptions → `uncaught_exception()`

```
struct task {  
    struct promise_type {  
        task get_return_object();  
        Awaitable initial_suspend();  
        void return_void();  
        void uncaught_exception();  
    };  
};
```

# Coroutines in C++ - promise

Finally, `final_suspend()` is called and `co_awaited`

```
struct task {  
    struct promise_type {  
        task get_return_object();  
        Awaitable initial_suspend();  
        void return_void();  
        void uncaught_exception();  
        Awaitable final_suspend() noexcept;  
    };  
};
```



# Coroutines in C++

We know everything to implement the idea!

Questions?

# Implementing the Idea

# Implementing the Idea

We learned how to

- suspend execution
- resume execution

# Implementing the Idea

1. Transform the loop body in a coroutine
2. Use `await` to interrupt execution, store the argument
3. Execute the batch when ready
4. Resume execution with the result

# Implementing the Idea

## Before

```
for(const User& user : users) {  
    const UserPrefs prefs = getUserPrefs({user.id}).at(0);  
    if (prefs.wantsEmailNotification) {  
        sendNotifs({prefs.notificationEmail});  
    }  
}
```

## After

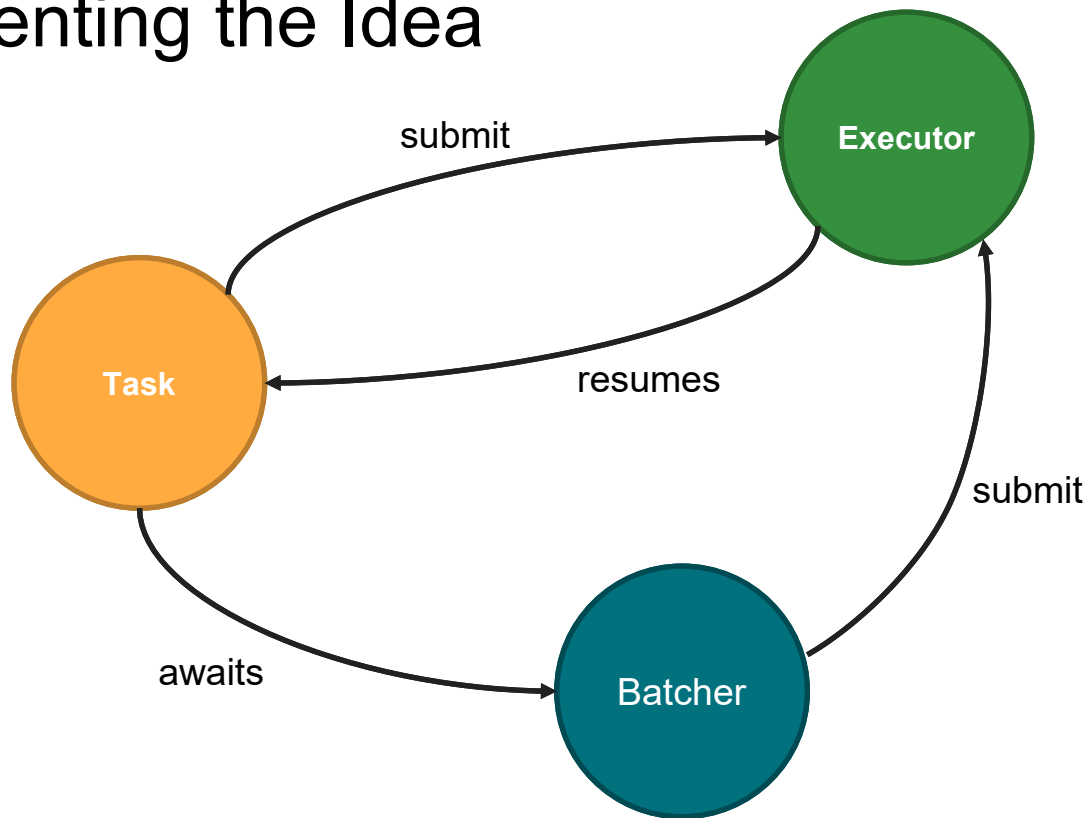
```
task sendEmail(const User& user) {  
    const UserPrefs prefs = co_await getUserPrefs(user.id);  
    if (prefs.wantsEmailNotification) {  
        co_await sendNotifs(prefs.notificationEmail);  
    }  
}
```

# Implementing the Idea

We'll need

1. a task
2. an executor
3. a batcher

# Implementing the Idea





# Implementing the Idea

```
Executor e;  
Batcher<...> getUserPrefs{...};  
Batcher<...> sendNotifs{...};  
  
auto sendEmail = [&](const User& user) -> task {  
    ...  
};  
  
for(auto& user : users) {  
    e.submit(sendEmail(user));  
}  
  
e.run_available();
```

# Implementing the Idea - Task

Let's look at the task

```
struct task {
```

```
};
```

# Implementing the Idea - Task

It must have a `promise_type`

```
struct task {  
    struct promise_type { ... };  
    using Handle = coroutine_handle<promise_type>  
  
};
```

# Implementing the Idea - Task

Similar to a `unique_ptr`

Owens the `coroutine_handle`

```
struct task {  
    struct promise_type { ... };  
    using Handle = coroutine_handle<promise_type>  
    explicit task(Handle ptr) : ptr_(ptr) {}  
    ...  
  
private:  
    Handle ptr_;  
};
```

# Implementing the Idea - Task

It's only moveable

```
struct task {  
    ...  
  
    task(const task&) = delete;  
    task(task&& t) noexcept  
        : ptr_(exchange(t.ptr_, nullptr)) {}  
  
    ...  
};
```

# Implementing the Idea - Task

Has RAI semantic

```
struct task {  
    ...  
  
    ~task() {  
        if (ptr_) ptr_.destroy();  
    }  
  
    ...  
};
```

# Implementing the Idea - Task

Exposes the `coroutine_handle`

```
struct task {  
    ...  
  
    Handle release() && {  
        return exchange(ptr_, nullptr);  
    }  
  
};
```

## Implementing the Idea - Task

Let's look at the `promise_type`

```
struct task::promise_type {  
  
  
  
  
  
  
  
  
};
```



# Implementing the Idea - Task

We immediately return when the coroutine is called

```
struct task::promise_type {  
    suspend_always initial_suspend() { return {}; }  
  
};
```

# Implementing the Idea - Task

We don't suspend when the coroutine terminates

```
struct task::promise_type {  
    suspend_always initial_suspend() { return {}; }  
    suspend_never final_suspend() noexcept { return {}; }  
  
};
```

# Implementing the Idea - Task

Ignore exceptions

```
struct task::promise_type {  
    suspend_always initial_suspend() { return {}; }  
    suspend_never final_suspend() noexcept { return {}; }  
    void unhandled_exception() { terminate(); }  
  
};
```

# Implementing the Idea - Task

We do nothing when we return

```
struct task::promise_type {  
    suspend_always initial_suspend() { return {}; }  
    suspend_never final_suspend() noexcept { return {}; }  
    void unhandled_exception() { terminate(); }  
    void return_void() {}  
};
```

# Implementing the Idea - Task

Let's look at how to implement `get_return_object()`

```
struct task::promise_type {  
    suspend_always initial_suspend() { return {}; }  
    suspend_never final_suspend() noexcept { return {}; }  
    void unhandled_exception() { terminate(); }  
    void return_void() {}  
    task get_return_object();  
};
```

# Implementing the Idea - Task

A reminder

```
struct std::coroutine_handle<promise_type> {  
    void resume();  
    void destroy();  
    promise_type & promise();  
    static coroutine_handle<promise_type> from_promise(promise_type&);  
};
```

# Implementing the Idea - Task

```
task get_return_object() {  
    Handle handle = Handle::from_promise(*this);  
    return task(handle);  
}
```

# Implementing the Idea - Task

Our `task` is complete!

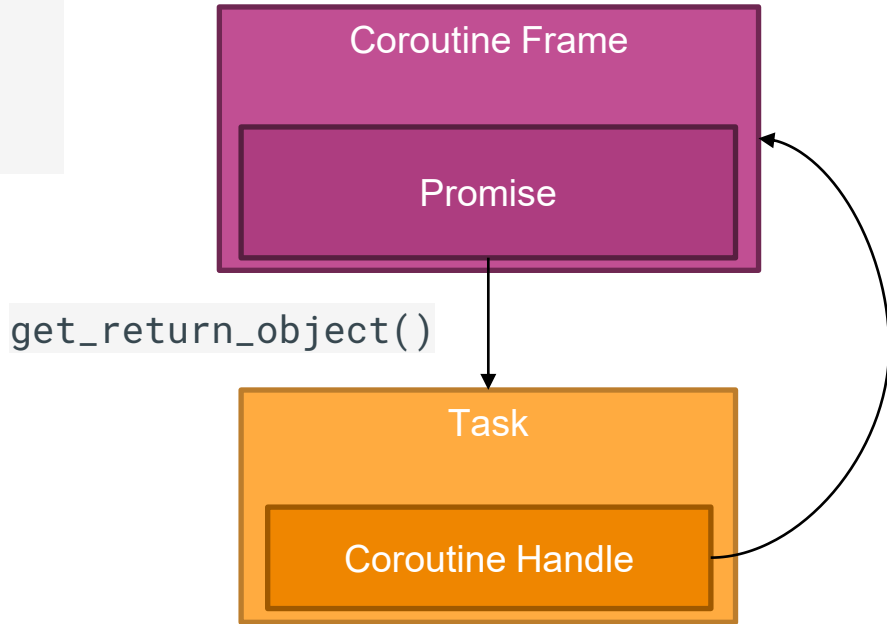
We can now write

```
task sendEmail(const User& user) {  
  
    co_return;  
}
```



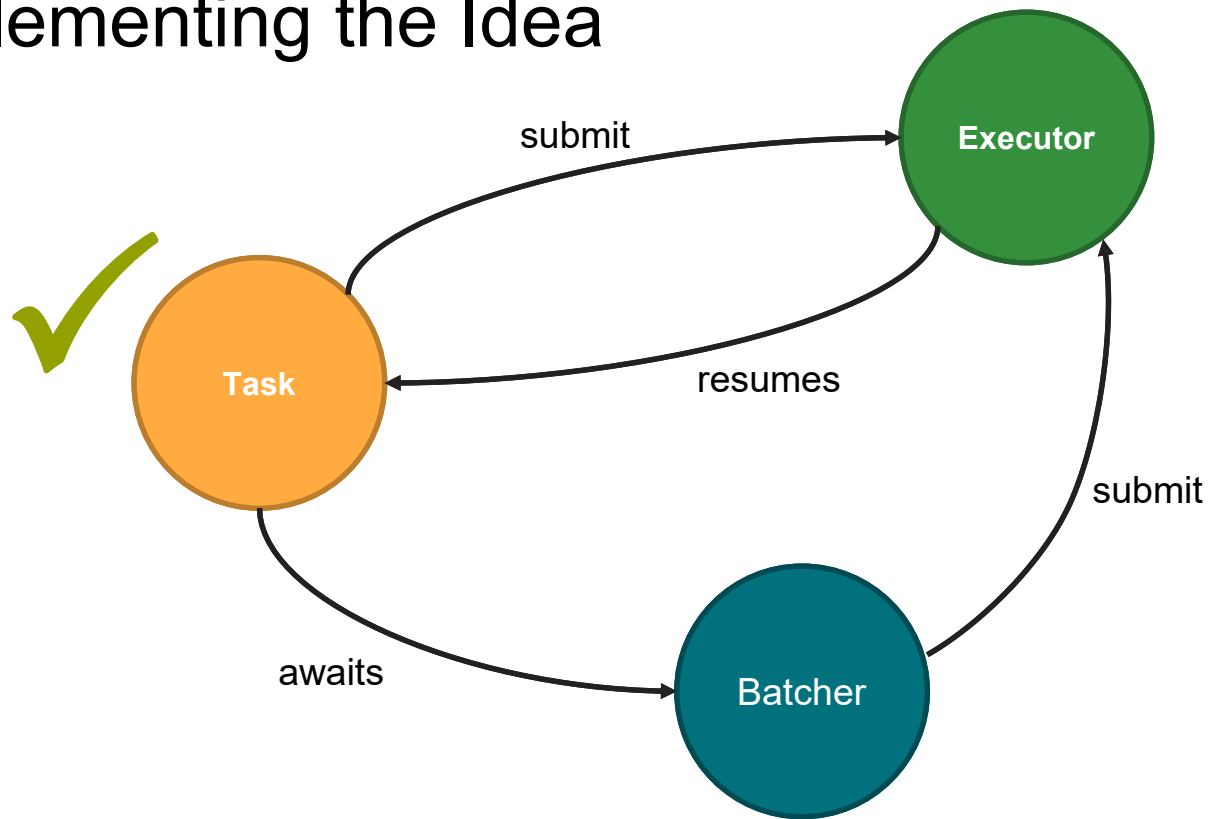
# Implementing the Idea - Task

```
task coro = sendEmail(user);
```



Questions?

# Implementing the Idea



# Implementing the Idea - Executor

We need something to execute the tasks

# Implementing the Idea - Executor

We need something to execute the `tasks`

Let's create an `Executor`

```
struct Executor {
```

```
};
```

# Implementing the Idea - Executor

## Supports

- submit coroutines
- execute pending coroutines

## Implementing the Idea - Executor

## Stores pending `coroutine_handles`

```
struct Executor {  
  
    deque<coroutine_handle<>> pending_  
};
```

## Implementing the Idea - Executor

## Submitting coroutines

```
struct Executor {  
  
    deque<coroutine_handle<>> pending_  
};
```



# Implementing the Idea - Executor

Can submit a task

```
struct Executor {  
    void submit(task t) {  
        auto handle = move(t).release();  
        pending_.push_back(handle);  
    }  
  
    deque<coroutine_handle<>> pending_;  
};
```

# Implementing the Idea - Executor

And arbitrary coroutines

```
struct Executor {  
    void submit(task t) {  
        auto handle = move(t).release();  
        pending_.push_back(handle);  
    }  
    void submit(vector<coroutine_handle<>> coros) {  
        // insert at the end  
    }  
  
    deque<coroutine_handle<>> pending_;  
};
```

## Implementing the Idea - Executor

## Executing coroutines

```
struct Executor {  
  
    deque<coroutine_handle<>> pending_  
};
```

# Implementing the Idea - Executor

Pull the first coroutine

```
struct Executor {  
    optional<coroutine_handle<>> pop_next_coro() {  
        if( pending_.empty() ) return nullopt;  
        auto first = pending_.front();  
        pending_.pop_front();  
        return first;  
    }  
  
    deque<coroutine_handle<>> pending_;  
};
```

# Implementing the Idea - Executor

Pull the first coroutine

```
struct Executor {  
    optional<coroutine_handle<>> pop_next_coro();  
  
    void run_available() {  
        for(auto next = pop_next_coro(); next; next = pop_next_coro()) {  
            next->resume();  
        }  
    }  
  
    deque<coroutine_handle<>> pending_;  
};
```

# Implementing the Idea - Executor

## Summary

```
struct Executor {  
    void submit(task t);  
    void submit(vector<coroutine_handle<>> coros);  
  
    optional<coroutine_handle<>> pop_next_coro();  
    void run_available();  
};
```

# Implementing the Idea - Executor


Our `Executor` is complete!

We can now write

```
Executor e;  
e.submit(sendEmail(user));  
e.run_available();
```

# Implementing the Idea - Executor

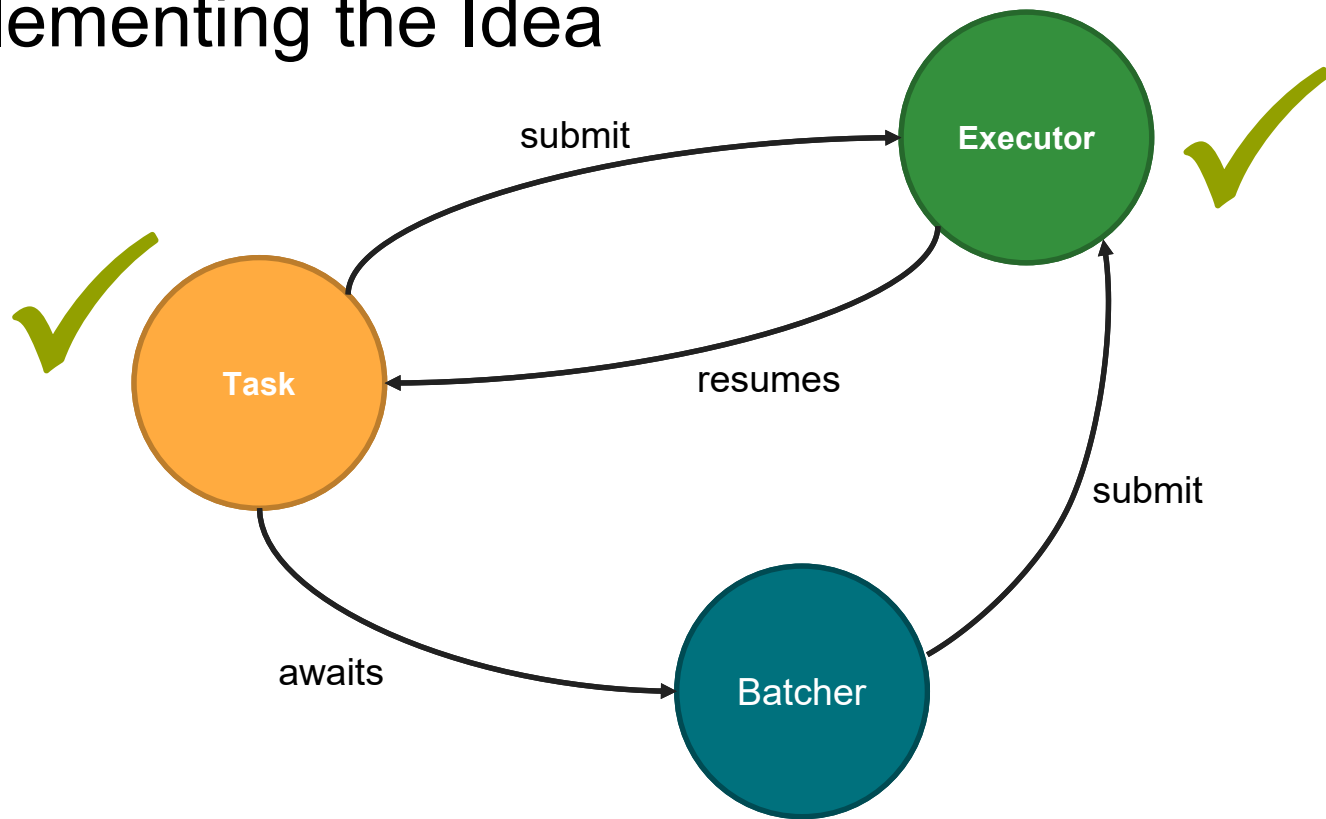
```
task coro = sendEmail(user);  
Executor e;  
e.submit(move(coro));  
e.run_available();
```





Questions?

# Implementing the Idea



# Implementing the Idea - Batcher

Let's support `co_await` now!

```
task sendEmail(const User& user) {  
    const UserPrefs prefs = co_await getUserPrefs(user.id);  
  
}
```

# Implementing the Idea - Batcher

`co_await` must be called on an `awaitable`

`getUserPrefs(user.id)` must return an `awaitable`




```
task sendEmail(const User& user) {  
    const UserPrefs prefs = co_await getUserPrefs(user.id);  
  
}
```

# Implementing the Idea - Batcher

Let's create a wrapper type: `Batcher`

- intercepts the calls to `getUserPrefs(user.id)`
- stores the arguments
- executes the function
- resumes the coroutines

# Implementing the Idea - Batcher





# Implementing the Idea - Batcher

Stores

- the executor

```
template<class T, class R>
struct Batchter {
    ...
    Executor& ex_;
};
```



# Implementing the Idea - Batcher

## Stores

- the executor
- the function

```
template<class T, class R>
struct Batcher {
    ...
    Executor& ex_;
    function<vector<R>(vector<T>)> op_;

};
```

# Implementing the Idea - Batcher

## Stores

- the executor
- the function and whether to execute it

```
template<class T, class R>
struct Batcher {
    ...
    Executor& ex_;
    function<vector<R>(vector<T>)> op_;
    function<bool(const vector<T>&)> should_exec_;

};
```

# Implementing the Idea - Batcher

## Stores

- the executor
- the function and whether to execute it
- the current batch

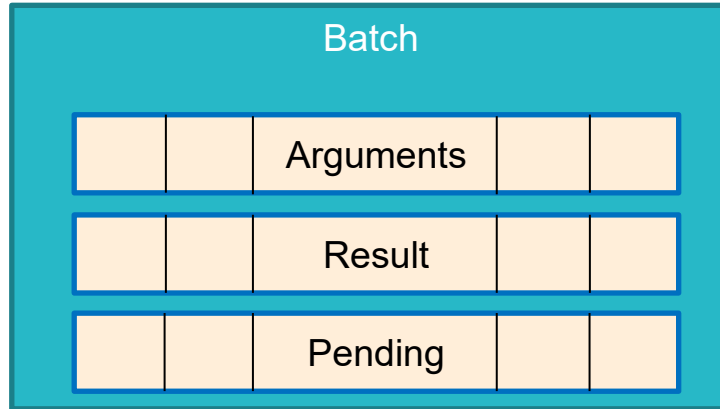
```
template<class T, class R>
struct Batcher {
    ...
    Executor& ex_;
    function<vector<R>(vector<T>)> op_;
    function<bool(const vector<T>&)> should_exec_;
    struct Batch { ... };
    shared_ptr<Batch> current_batch_;
};
```

## Implementing the Idea - Batcher

## Let's look at the Batch

```
struct Batcher::Batch {  
  
  
  
  
  
  
};
```

# Implementing the Idea - Batcher



# Implementing the Idea - Batcher

```
struct Batcher::Batch {  
    vector<T> arguments;  
    vector<R> result;  
    vector<coroutine_handle<>> pending;  
};
```

# Implementing the Idea - Batcher


`Batcher` is callable

and returns an awaitable object

```
template<class T, class R>
struct Batcher {
    struct Awaitable { ... };

    Awaitable operator()(T arg);
    ...
};
```

# Implementing the Idea - Batcher





# Implementing the Idea - Batcher

```
struct Batcher::Awaitable {  
  
    Batcher& batcher_;  
    shared_ptr<Batch> batch_;  
    size_t index_ = -1;  
};
```

# Implementing the Idea - Batcher

How is the `Awaitable` constructed?

```
template<class T, class R>
struct Batcher {

    Awaitable operator()(T arg) {
        current_batch_>args.push_back(arg);
        size_t index = current_batch_>args.size() - 1;
        return Awaitable{*this, current_batch_, index};
    }
    ...
};
```

# Implementing the Idea - Batcher

`Awaitable` needs to implement the special methods

Let's do that!

# Implementing the Idea - Batcher

Check if it's ready

```
struct Batcher::Awaitable {  
    bool await_ready() {  
        return batcher_.maybe_execute();  
    }  
    ...  
};
```

# Implementing the Idea - Batcher

Get the result once ready

```
struct Batcher::Awaitable {  
    bool await_ready();  
    R await_resume() {  
        return batch_>results.at(index_);  
    }  
    ...  
};
```

# Implementing the Idea - Batcher

## Suspend

```
struct Batcher::Awaitable {  
    bool await_ready();  
    R await_resume();  
    void await_suspend(coroutine_handle<> h) {  
        batch_->pending_.push_back(h);  
    }  
    ...  
};
```

# Implementing the Idea - Batcher

## Suspend

```
struct Batcher::Awaitable {
    bool await_ready();
    R await_resume();
    coroutine_handle<> await_suspend(coroutine_handle<> h) {
        batch_->pending_.push_back(h);
        auto maybe_coro = batch_.executor_.pop_next_coro();
        return maybe_coro.value_or(std::noop_coroutine());
    }
    ...
};
```

# Implementing the Idea - Batcher

Only thing left: `maybe_execute()`

```
template<class T, class R>
struct Batcher {
    bool maybe_execute() {
        if( !should_exec_(batch_->args) ) return false;
        batch_->results = op_(batch_->args);
        executor_.submit(batch_->pending);
        batch_ = make_shared<Batch>();
        return true;
    }
};
```






# Implementing the Idea - Batcher

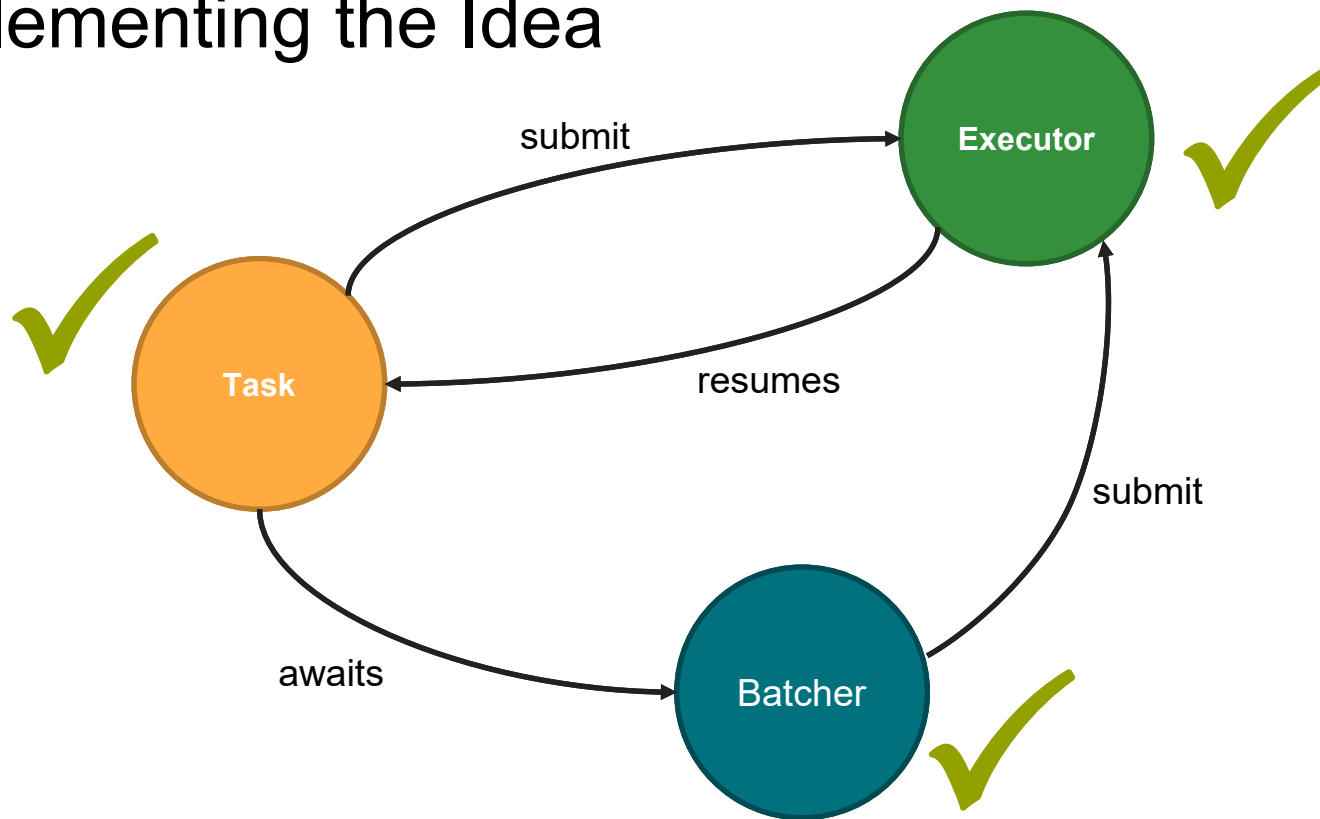
Batcher is complete!

```
Executor e;  
Batcher<User, UserPref> getPrefs{e, getUserPrefs, size_eq(10)};  
  
auto sendEmail = [&](const User& user) -> task {  
    const UserPrefs prefs = co_await getPrefs(user.id);  
  
};
```

# Implementing the Idea - Batcher



# Implementing the Idea



Questions?

# Putting all together

# Putting all together

```
Executor e;  
Batcher<...> getUserPrefs{...};  
Batcher<...> sendNotifs{...};  
  
auto sendEmail = [&](const User& user) -> task {  
    const UserPrefs prefs = co_await getUserPrefs(user.id);  
    if (prefs.wantsEmailNotification) {  
        co_await sendNotifs(prefs.notificationEmail);  
    }  
};  
  
for(auto& user : users) {  
    e.submit(sendEmail(user));  
}  
  
e.run_available();
```




# Problem

What if we finish the available coroutines before completing?

Example

- Batching with no limit
- Batching limit not a perfect divisor of the initial vector

# Problem





# Problem

## Result

- the executor has no pending coroutines
- the batch has pending coroutines

Some tasks do not complete

# Solution

Force the batch execution

```
template<class T, class R>
struct Batcher {
    bool maybe_execute(bool force = false) {
        bool exec = force || should_exec_(batch_->args);
        if( !exec ) return false;
        batch_->results = op_(batch_->args);
        executor_.submit(batch_->pending);
        batch_ = make_shared<Batch>();
        return true;
    }
};
```

# Problem

When do we stop calling force?

# Solution

When do we stop calling force?

When all tasks completed

# Solution

Let's keep track of the tasks

```
struct Executor {  
    void submit(task t) {  
        counter_++;  
        auto handle = move(t).release();  
        handle.promise().counter_ = &counter_;  
        pending_.push_back(handle);  
    }  
  
    size_t counter_ = 0;  
};
```

# Solution

Decrease the counter when a task completes

```
struct task::promise_type {  
    void return_void() {  
        *counter -= 1;  
    }  
    size_t* counter_;  
};
```

# Solution

Return whether all tasks completed

```
struct Executor {  
    bool run_available() {  
        for(auto next = pop_next_coro(); next; next = pop_next_coro()) {  
            next.resume();  
        }  
        return counter_ == 0;  
    }  
};
```

# Solution

```
Executor e;  
Batcher<...> getUserPrefs{...};  
Batcher<...> sendNotifs{...};  
  
auto sendEmail = [&](const User& user) -> task {...};  
  
for(auto& user : users) {  
    e.submit(sendEmail(user));  
}  
  
while(e.run_available()) {  
    getUserPrefs.maybe_execute(true);  
    sendNotifs.maybe_execute(true);  
}
```



# Problem

We are calling `maybe_execute()` on every batcher

It's unnecessary

# Solution

Execute only the first `maybe_execute()`

# Solution

Skip the execution if no args




```
template<class T, class R>
struct Batcher {
    bool maybe_execute(bool force = false) {
        bool exec = force || should_exec_(batch_>args);
        if( !exec || batch_>args.empty() ) return false;
        batch_>results = op_(batch_>args);
        executor_.submit(batch_>pending);
        batch_ = make_shared<Batch>();
        return true;
    }
};
```

# Solution

Execute only the first `Batcher`

```
...  
while(e.run_available()) {  
    getUserPrefs.maybe_execute(true) ||  
    sendNotifs.maybe_execute(true);  
}
```

# Problem



# Final Solution

```
Executor e;  
Batcher<...> getUserPrefs{...};  
Batcher<...> sendNotifs{...};  
  
auto sendEmail = [&](const User& user) -> task {  
    const UserPrefs prefs = co_await getUserPrefs(user.id);  
    if (prefs.wantsEmailNotification) {  
        co_await sendNotifs(prefs.notificationEmail);  
    }  
};  
  
for(auto& user : users) { e.submit(sendEmail(user)); }  
  
run_to_completion(e, getUserPrefs, sendNotifs);
```

Questions?

# Benchmark



# Benchmark

We implemented the `task`, an `executor` and a `Batcher` with coroutines

We can now automatically batch our operations

How does it compare to manual batching?

# Manual Batching

```
for(const User& user : users) {  
    const UserPrefs prefs =  
        getUserPrefs({user.id}).at(0);  
    if (prefs.wantsEmailNotification) {  
        sendNotifs({prefs.notificationEmail});  
    }  
}
```

```
vector<UserId> userIds;  
transform(users.begin(), users.end(),  
    back_inserter(userIds),  
    [](auto& user) { return user.id; });
```

```
vector<UserPrefs> preferences =  
    getUserPrefs(userIds);  
remove_if(preferences.begin(),  
    preference.end(),  
    [](auto& pref) {  
        return !pref.wantsEmailNotification;});
```

```
vector<EmailAddress> emails;  
transform(preferences.begin(),  
    preference.end(),  
    back_inserter(emails),  
    [](auto& pref) {  
        return pref.notificationEmail; });  
  
sendNotification(emails);
```

# Coro Batching

```
for(const User& user : users) {  
    const UserPrefs prefs =  
        getUserPrefs({user.id}).at(0);  
    if (prefs.wantsEmailNotification) {  
        sendNotifs({prefs.notificationEmail});  
    }  
}
```

```
Executor e;  
Batcher<...> getUserPrefs{...};  
Batcher<...> sendNotifs{...};
```

```
auto sendEmail = [&](const User& user)->task {  
    const UserPrefs prefs =  
        co_await getUserPrefs(user.id);  
    if (prefs.wantsEmailNotification) {  
        co_await sendNotifs(  
            prefs.notificationEmail);  
    }  
};
```

```
for(auto& user : users) {  
    e.submit(sendEmail(user));  
}
```

```
run_to_completion(e, getUserPrefs, sendNotifs);
```

# Benchmark

Is performance comparable?

Simulation

- 10 ms delay per call
- 50 us delay per item processed (100 users, 50 email sent)

Test	Time (us)	Iterations	Baseline
No Batching	1'552'542	13	49.99
Manual Batching	31'055	492	-
Coro Batching	27'748	564	0.89

# Benchmark

Overhead?

Simulation without delays

Test	Time (ns)	Iterations	Baseline
No Batching	17'004	446'422	1.61
Manual Batching	10'586	655'351	-
Coro Batching	30'490	224'300	2.88

# Benchmark

Overhead?

Simulation without delays

Test	Time (ns)	Iterations	Baseline
No Batching	17'004	446'422	1.61
Manual Batching	10'586	655'351	-
Coro Batching	30'490	224'300	2.88
Coro Batching (Opt)	12'546	548'914	1.18

# Conclusion

# Conclusion

A new point on the performance  $\leftrightarrow$  readability trade-off

## Learned

- Main customization points of coroutines
- Coroutines can be used for more than just async programming
- Coroutines are powerful, but at times complicated



# Conclusion

Find the full implementation at

[github.com/MakersF/cppcon-2021-corobatch](https://github.com/MakersF/cppcon-2021-corobatch)

Library exploring the concept

[github.com/MakersF/corobatch](https://github.com/MakersF/corobatch)

Thank you!

# Useful Material

- Lewis Baker – [lewissbaker.github.io/](https://lewissbaker.github.io/)
- Dawid Pilarski – [blog.panicsoftware.com](https://blog.panicsoftware.com)
- Gor Nishanov – [“Nano-coroutines to the Rescue!”](#)